

Analysis of semaphore implementations

by Alexei Finski

The following implementation is provided at

<https://sites.cs.ucsb.edu/~rich/class/cs170/notes/Semaphores/index.html>:

```
void P(sema *s){
    pthread_mutex_lock(&s->lock);
    s->value--;
    while (s->value < 0){
        if (s->waiters < -1 * s->value){
            s->waiters++;
            pthread_cond_wait(&s->wait, &s->lock);
            s->waiters--;
        }else{
            break;
        }
    }
    pthread_mutex_unlock(&s->lock);
    return;
}

void V(sema *s){
    pthread_mutex_lock(&s->lock);
    s->value++;
    if (s->value <= 0){
        pthread_cond_signal(&s->wait);
    }
    pthread_mutex_unlock(&s->lock);
}
```

`s->waiters` is initialized to 0 and is non-negative because for every decrement operation there is a preceding increment operation. Thus, for any `s->value` ≥ 0 `s->waiters` $< -1 * s->value$ is false. Based on this observation, the above implementation simplifies to the following implementation of P, which will be considered.

```
void P(sema *s){
    pthread_mutex_lock(&s->lock);
    s->value--;
    while (s->waiters < -1 * s->value){
        s->waiters++;
        pthread_cond_wait(&s->wait, &s->lock);
        s->waiters--;
    }
    pthread_mutex_unlock(&s->lock);
}
```

Wlog, let `s->value` == -2 and `s->waiters` == 2. Consider the following example:

- 1) Thread A calls V, increments `s->value` to -1, and signals with `pthread_cond_signal`.
- 2) Thread A continues running, calls P, decrements `s->value` to -2, and does not call `pthread_cond_wait`.
- 3) Thread B is awakened by the signal from `pthread_cond_signal` called by thread A, reacquires mutex, and decrements `s->waiters` to 1. Because `s->value` == -2 and `s->waiters` == 1, thread B does not exit the while loop, increments `s->waiters` to 2 and calls `pthread_cond_wait`.

Thus, thread A “received” its own signal and avoided being blocked, whereas thread B continues being blocked. Therefore, the above implementation can lead to thread starvation.

The problem is addressed by guaranteeing a call to *pthread_cond_wait* (*cond_wait*) with a do...while loop for each thread that calls P (*sem_wait*) when the value of a semaphore is ≤ 0 after mutex acquisition in P (*sem_wait*), as provided in The Little Book of Semaphores by Allen B. Downey (Version 2.2.1):

```
void sem_wait(Semaphore *semaphore){
    mutex_lock(semaphore->mutex);
    semaphore->value--;
    if (semaphore->value < 0){
        do{
            cond_wait(semaphore->cond, semaphore->mutex);
        }while (semaphore->wakeups < 1);
        semaphore->wakeups--;
    }
    mutex_unlock(semaphore->mutex);
}
```

```
void sem_signal(Semaphore *semaphore){
    mutex_lock(semaphore->mutex);
    semaphore->value++;
    if (semaphore->value <= 0){
        semaphore->wakeups++;
        cond_signal(semaphore->cond);
    }
    mutex_unlock(semaphore->mutex);
}
```